

See discussions, stats, and author profiles for this publication at: <https://www.researchgate.net/publication/240648947>

Topology cuts: A novel min-cut/max-flow algorithm for topology preserving segmentation in ND images

Article in *Computer Vision and Image Understanding* · October 2008

DOI: 10.1016/j.cviu.2008.07.008 · Source: DBLP

CITATIONS

79

READS

262

4 authors:



[Yun Zeng](#)

Harvard University

18 PUBLICATIONS 454 CITATIONS

[SEE PROFILE](#)



[Dimitris Samaras](#)

Stony Brook University

229 PUBLICATIONS 6,793 CITATIONS

[SEE PROFILE](#)



[Wei Chen](#)

95 PUBLICATIONS 1,514 CITATIONS

[SEE PROFILE](#)



[Qunsheng Peng](#)

Zhejiang University

328 PUBLICATIONS 3,903 CITATIONS

[SEE PROFILE](#)

Some of the authors of this publication are also working on these related projects:



3D Shape Matching [View project](#)



Visual Analytics [View project](#)



Topology cuts: A novel min-cut/max-flow algorithm for topology preserving segmentation in N-D images

Yun Zeng^{a,*}, Dimitris Samaras^a, Wei Chen^b, Qunsheng Peng^b

^a Computer Science Department, Stony Brook University, Stony Brook, NY 11794-4400, USA

^b State Key Lab of CAD&CG, Zhejiang University, Hangzhou 310028, PR China

ARTICLE INFO

Article history:

Received 1 November 2007

Accepted 10 July 2008

Available online 5 August 2008

Keywords:

Image segmentation
Min-cut/max-flow
Topology preservation
Topology cuts
Graph cuts

ABSTRACT

Topology is an important prior in many image segmentation tasks. In this paper, we design and implement a novel graph-based min-cut/max-flow algorithm that incorporates topology priors as global constraints. We show that the optimization of the energy function we consider here is NP-hard. However, our algorithm is guaranteed to find an approximate solution that conforms to the initialization, which is a desirable property in many applications since the globally optimum solution does not consider any initialization information. The key innovation of our algorithm is the organization of the search for maximum flow in a way that allows consideration of topology constraints. In order to achieve this, we introduce a label attribute for each node to explicitly handle the topology constraints, and we use a distance map to keep track of those nodes that are closest to the boundary. We employ the bucket priority queue data structure that records nodes of equal distance and we efficiently extract the node with minimal distance value. Our methodology of embedding distance functions in a graph-based algorithm is general and can also account for other geometric priors. Experimental results show that our algorithm can efficiently handle segmentation cases that are challenging for graph-cut algorithms. Furthermore, our algorithm is a natural choice for problems with rich topology priors such as object tracking.

© 2008 Elsevier Inc. All rights reserved.

1. Introduction

Many computer vision problems such as segmentation, stereo reconstruction and image restoration, can be formulated as a minimization of an energy function [26]. These energy functions are naturally divided into two groups: continuous and discrete. For the problem of image segmentation, the level sets method [25] is a representative model in the continuous community, while the graph-based Markov Random Fields (MRFs) [13] is a very popular model in the discrete group. One very efficient algorithm for solving a subclass of the MRF energy function is the graph cuts algorithm [8]. In recent years, there has been a number of works showing the close relationship between level sets and graph cuts ([4,6,19], etc.), and how shape priors can be incorporated into the graph cuts framework ([12,22]). In this work, we show how the idea of topology preserving segmentation from the level sets literature [14] can be transposed to the graph-based algorithms. We propose the first min-cut/max-flow algorithm that is designed to explicitly incorporate topology as a global constraint in the segmentation. We call our new algorithm Topology Cuts, in analogy to the popular graph cuts algorithm [8].

Topology as a prior is available in many applications. For example, the anatomy of human tissues provides important topological constraints that ensure the correctness in biomedical image segmentation. Existing techniques that enforce topology constraints into the graph cuts algorithm, do so by simply tuning the parameters of the energy function [3,7]. This scheme usually requires intense user interactions and is not applicable in cases where user manipulation is difficult. In contrast, we propose to embed the topological constraint into the discrete min-cut/max-flow algorithm, which leads to a new and efficient way of considering global topology information for the general problem of topology preservation.

Our work is inspired by the topology preserving level set method of [14]. This algorithm makes use of digital topology theory for N-D images [2] to detect topology changes during the evolution of level sets. Taking advantage of the fact that the level set functions are solved in a gradient descent manner, and assuming that the change of sign for the pixels only occurs one pixel at a time on the boundary of the evolving objects, the topology of the object can be easily controlled. The advantage of our method over [14] is the speed-up and numerical stability inherent to discrete max-flow methods. In addition, our method has a guaranteed convergence property.

However, transferring the idea of topology preserving evolution from the continuous level sets algorithm to the discrete

* Corresponding author.

E-mail address: yzeng@cs.sunysb.edu (Y. Zeng).

graph-based algorithm is not straightforward. The main difficulty lies in the fact that previous graph-cut implementations [1,5,15] are inherently topology-free and thus not conducive to topology considerations during the search of max-flow. To make these considerations possible for the discrete graph-based algorithm, we introduce the following elements.

- (1) An *F/B* label attribute is introduced to explicitly handle the topology property in the image. This resolves an ambiguity in the existing graph cuts algorithms, i.e., it is possible that the labels for a subset of the graph's nodes can be changed without changing the optimal solution (multiple solutions for the energy minimization problem). Existing algorithms set these nodes' labels to a default label, which unavoidably leads to topological errors.
- (2) An *initialization* step is used to provide the graph with initial topology information.
- (3) The computation of max-flow is divided into *inter-label* and *intra-label* stages, to facilitate the propagation of topology information during the search for the minimum of the energy function.
- (4) A *distance map* (function) which keeps track of the nodes that are closest to the current boundary between the different label sets is set in the beginning and is updated during the computation.
- (5) To efficiently insert and extract nodes on the current evolving boundary (the level set of the distance map), we use the *bucket priority queue* data structure [9,11], which only requires time of $O(1)$ complexity for each insertion and extraction operation. Hence, there is no loss of efficiency compared to the previous graph-cut algorithms. Our algorithm shares the same complexity with the widely used graph-cut implementation [5], and in practice it runs in comparable speed.

The contributions of this paper can be summarized as follows:

- To the best of our knowledge, this is the first work that incorporates a global topology prior into the design of the discrete graph-based min-cut/max-flow algorithms.
- We prove that enforcing global optimality of the solution while considering topology constraints is NP-hard. That means any algorithm that enforces topology constraints either interactively or automatically can not obtain the global optimum.
- In the design of our algorithm, we combine concepts from the level-set literature (such as distance maps and level-set evolution [24,25]) into an efficient discrete graph-based algorithm. The techniques we use here are general and define a new way of incorporating geometric prior knowledge into the existing graph-cut models/algorithms such as curvature or shape priors.

Additionally, our new algorithm is suitable for the concept of multilevel banded graph cuts [23] to fairly speedup the computation. In experiments, we show that our algorithm achieves more meaningful and visually better results compared with graph cuts for problems where topology information is available, e.g., image segmentation and object tracking.

1.1. Organization of this paper

In Section 2, we review the essential background for describing our new algorithm. Section 3 discusses the primal–dual schema in the min-cut/max-flow algorithm and the overall principles of our algorithm. Section 4 gives the formulation of the topology-cut

problem and proves its NP-hardness. Section 5 explains the design of our new algorithm. Section 6 discusses the detailed implementation. Section 7 analyzes our algorithm with respect to convergence, topology preservation and time complexity. The experimental results are presented in Section 8. Finally, we conclude and outline future work in Section 9.

2. Preliminaries

2.1. Digital topology

Here, we discuss two key concepts in the topology of digital images[2]: connectivity and simple point.

The *connectivity* of a digital image specifies the condition of adjacency that two points must fulfill, in order for the foreground *F* and the background *B* to be considered connected respectively.

To ensure this (Fig. 1), different connectivity for the object and the background must be specified. For 2D images, valid (foreground, background) connectivity pairs are (4,8) and (8,4); for 3D images, the valid pairs are (6,18), (6,26), (18,6) and (26,6).

With the clarification of connectivity in the digital image discussed above, a *simple point* is defined as a point whose change from foreground to background or vice versa, does not change the number of connected components of both the foreground and background. A simple point can be efficiently computed using the concept of topological number [2]. The definition of simple point has been relaxed in [29,32]. In this paper, we adopt the basic definition of simple points discussed above for detecting topology change. With the above definitions, we define that two binary images are topologically equivalent if one image can be changed to another by updating only simple points. Note that this equivalence relation is symmetric.

2.2. Energy minimization and min-cut/max-flow

The MRF energy function [13] solved by graph cuts can be formulated as:

$$\inf_{x_p, p \in \mathcal{V}} \left\{ \sum_{p \in \mathcal{V}} D(x_p) + \sum_{(p,q) \in \mathcal{E}} V_{pq}(x_p, x_q) \right\}, \quad \text{with } x_p \in \{0, 1\}, p \in \mathcal{V}. \quad (1)$$

Here, \mathcal{V} and \mathcal{E} usually denote the image pixels and their pairwise relationships, respectively. If the terms of the MRF energy function have the form $D(x_p) = D_p^0(1 - x_p) + D_p^1 x_p$ and $V_{pq}(x_p, x_q) = w_{pq}((1 - x_p)x_q + (1 - x_q)x_p)$ with $w_{pq} \geq 0$ (submodular condition), we can define a graph \mathcal{G} with a source terminal *s*, a sink




Fig. 1. The black grids denote the foreground object and the white grids denote the background. (a) Illustrates an example of topological paradox. (b) A valid connectivity pair.

terminal t , and nodes $\{p|p \in \mathcal{V}\}$. The capacity from s to each node p is defined as D_p^s , the capacity from each p to t is defined as D_p^t , the capacity between neighboring nodes p, q is defined as w_{pq} .

Adopting the notation in [5], a s/t cut of a graph \mathcal{G} is a partition of the nodes and terminals into two disjoint subsets S and T with $s \in S$ and $t \in T$. Also, the cost of a cut $C = \{S, T\}$ is the sum of the costs of all the edges (p, q) where $p \in S$ and $q \in T$. It has been proven [1] that the optimal solution of the binary energy function (1) corresponds to a *min-cut* of the graph \mathcal{G} , which is the minimal cut of all the cuts in the graph.

Finding the min-cut of a graph is equivalent to computing a maximum flow from s to t [1]. In general, algorithms for solving this min-cut/max-flow problem fall into two groups: augmenting path and push-relabel [1] techniques. Here we review a popular algorithm [5] based on augmenting paths that is closely related to our method.

2.3. A dynamic tree implementation of the s/t cut

Generally speaking, the idea of the augmenting path algorithm [1] is to iteratively search a non-saturated path from s to t and push the maximal possible flow along this path. When no more such paths can be found, the maximum flow has been reached.

There is a number of ways to search non-saturated paths between two terminals. The efficient algorithm in [5] searches non-saturated paths by growing two trees from both the source and the sink. This idea can be efficiently implemented using the dynamic tree data structure.

A dynamic tree grows by adding non-saturated edges dynamically. As Fig. 2 shows, two non-overlapping dynamic trees S and T are maintained. Each node either belongs to one of the two trees or is “free”. A node that belongs to a tree can be in either “active” or “passive” state. An active node is at the border of the tree while a passive node is inside the tree (See Fig. 2).

To find a non-saturated path in the graph, three stages are iteratively repeated:

- “growth”: grow trees S and T until they meet in the middle, giving a non-saturated $s \rightarrow t$ path,
- “augmentation”: push the maximum possible flow along this path, breaking the trees into a forest,
- “adoption”: restore the single tree structure of the two trees by finding a new parent for the isolated parts. If no such parent can be found, they become free nodes.

The algorithm stops when there are remaining active nodes in the two trees.

3. Tree membership and the primal–dual solution to the s/t cut

If we relax the variables of the discrete optimization problem (1) to be continuous, its duality can be formulated as in [30]




Fig. 2. The dynamic tree implementation of the s/t cut.




Fig. 3. A planar graph \mathcal{H} with its connected vertex covering (a), and its embedding into the grid image domain \mathcal{G} (b).

$$\begin{aligned} \max \quad & f_{ts} \\ \text{s.t.} \quad & f_{pq} \leq w_{pq}, (p, q) \in \mathcal{E} \\ & \sum_{p:(p,q) \in \mathcal{E}} f_{pq} - \sum_{p:(q,p) \in \mathcal{E}} f_{qp} \leq 0, \quad q \in \mathcal{V} \\ & f_{pq} \geq 0 (p, q) \in \mathcal{E}. \end{aligned} \quad (2)$$

It can be shown that for any feasible solution of Eq. (2), $f_{ts} \leq c_{\text{opt}}$ where c_{opt} is the optimal solution of the primal Eq. (1). Thus finding the max-flow of the graph corresponds to finding a lower bound of the primal problem. In the ideal situation this lower bound reaches the optimal solution of the primal Eq. (1). However, for many cases, e.g., optimizing energy of the primal Eq. (1) with additional constraints, this lower bound can not reach a global optimum. Thus it can only be used as guidance to the approximate solution of the primal problem in a similar manner as [16–18].

In the standard implementations of graph cuts, the label of each node (foreground or background) is normally determined by whether there is a non-saturated path from it to s or t when the max-flow is reached, namely, these algorithms determine the label of each node by its *tree membership*. Using tree membership to determine the label of each node makes it impossible to consider topology properties during the max-flow computation. This is because the tree membership of each node is updated in an irregular order and thus can not contain any topological information (graph cuts are inherently topology-free).

With the above discussion of the primal–dual relation, our new algorithm can be regarded as solving the 0/1 optimization in its dual space (namely, finding max-flow), while making sure that its intermediate primal solution, which is represented by tree-membership, conforms to topology constraint. Since our algorithm is solved in the dual domain, even without considering additional constraints, each update of the primal solution may not necessarily lead to a lower energy Eq. (1). Furthermore, similar to [14], the feasible solution domain of our problem is non-convex and therefore our method may obtain multiple segmentations depending on initialization. Even though, as shown in Section 4, we can not guarantee global optimality for any of these solutions, they all have the following properties. They are either (1) the original cut when the topology constraint is not violated, or (2) if the original cut violates the topology constraint, we select the segmentation that (a) respects the topology constraint and (b) is derived from the homogeneous propagation of the boundary of the initialization.

4. The topology cuts problem

The energy function that combines the MRF formulation and digital topology on the image grid is

$$\begin{aligned} \inf_{x_p, p \in \mathcal{V}} \quad & \left\{ \sum_{p \in \mathcal{V}} D(x_p) + \sum_{(p,q) \in \mathcal{E}} V_{pq}(x_p, x_q) \right\}, \\ \text{s.t.} \quad & \mathcal{T} = \mathcal{T}_{\text{init}} \\ & \text{with } x_p \in \{0, 1\}, p \in \mathcal{V}. \end{aligned} \quad (3)$$

Here, \mathcal{T} denotes the topology of the 0/1 labeled image as defined in [2]. $\mathcal{T}_{\text{init}}$ is the initial topology information that is assigned to the image either interactively, or automatically as shown in our tracking example. The meanings of the other notations are the same as in Eq. (1). Note that here we only consider the hard-constraints on the topology of the image. However, soft-constraints can be conveniently introduced by considering alternative definitions of simple points [29].

Theorem 1. *The topology cuts (TP-CUT) problem (3) is NP-hard.*

Proof. To prove the NP-hardness of the TP-CUT problem, we reduce from the connected vertex-covering (CVC) problem [28].

Definition 1. Connected vertex-covering problem Given a planar graph $\mathcal{H} = (V_h, E_h)$ with maximum degree 4 for each node and an integer $K (\geq 1)$, is there a connected subset $V' \subset V_h$ such that $|V'| \leq K$ and for each edge $e \in E_h$, at least one of its two end nodes is in V' (a vertex cover)?

The CVC problem has been proven to be NP-hard in [28,31]. To reduce from the CVC problem to our TP-CUT problem, we use the following techniques.

1. A planar graph with maximum degree 4 can be embedded into the image grid domain \mathcal{G} [28] (Fig. 3).
2. The weight of the edges for neighboring pixels is set to be zero, so we only have to consider the first term in (3).
3. The 0/1 label of a node on the image \mathcal{G} can be “fixed” by setting a sufficiently large weight to the edge which links it to the source (source edge) or the sink nodes (sink edge). This means if we are to change the label of the node, a very large penalty would be added to the energy function (3) (Fig. 4). We call such a large weight *infinite weight*.
4. We may give some nodes on the image the “freedom” to change their label by setting the weight of the source/sink edge to be a specifically designed value. The weight of the sink edges for the nodes corresponding to vertices in \mathcal{H} (vertex node) is set to be $|E_h| + 1$, where $|E_h|$ is the number of edges of \mathcal{H} (Fig. 4(b)). The weights of the sink edge of all the other nodes (grid) which do not lie on the embedded graph are set to be infinite (Fig. 4(c)).
5. To respect topology constraints, we pick one node for each embedded edge in \mathcal{G} and set the weight of its source edge to be a small value (Fig. 4(a)). Intuitively, this constructs a “door” on each edge to allow different closed regions in the original planar graph connected to each other (we call such nodes *door nodes*). Such door nodes can be labeled 1 only when the connectivity of the 0-labeled foreground would not be broken. In other words, using these nodes, it is always possible to connect an enclosed 1-labeled region (by a 0-labeled “wall”) to the outside

of the wall by opening one door on the wall without separating the wall into two parts. When such a door is opened, a small penalty is added to the total energy function (Fig. 4(a)).

6. With the above construction, a CVC problem has a vertex cover of size no greater than K if and only if its corresponding TP-CUT problem with the constraint that the all the foreground/background nodes are connected, has an optimal solution whose energy is less than $(K + 1)(|E_h| + 1)$. In this way a one-to-one correspondence between the CVC problem and the TP-CUT problem is established.

An example of the reduction is shown in Fig. 5.

The reduction is constructed as follows: (a) A planar graph \mathcal{H} is embedded into an image \mathcal{G} with two terminals (source/sink). (b) The weight of the edges for neighboring pixels is set to be zero. The other edges in \mathcal{G} are those edges linking the nodes to the two terminals (source/sink edges). (c) The label of nodes on the image that do not lie on the embedded graph \mathcal{H} is “fixed” to be background by setting a sufficiently large (“infinite”) weight to each node’s sink edge. (d) To respect topology constraints, we pick one node for each embedded edge and set the weight of its source edge to be one. We also fix the label of the other nodes on these edges to be foreground by setting an infinite weight for their source edges. (e) The weight of the sink edges for the nodes corresponding to vertices in \mathcal{H} is set to be $|E_h| + 1$, where $|E_h|$ is the number of edges of \mathcal{H} ; With the above construction, a CVC problem has a vertex cover of size no greater than K if and only if its corresponding TP-CUT problem with the constraint that all the foreground/background nodes are connected, has an optimal solution whose energy is less than $(K + 1)(|E_h| + 1)$.

Since the TP-CUT problem is NP-hard, our goal in this paper is to design an efficient algorithm that finds a local optimum. We update the binary partition of the image by solving the above energy function using the standard min-cut/max-flow algorithm while making sure that each update does not violate the topology constraint. Our algorithm is able to handle all the energy functions that can be solved by graph cuts [20] with an additional topology constraint. In the rest of this paper, we discuss our algorithm in detail.

5. Design of the topology cuts algorithm

In this section, we discuss the novel aspects of our algorithm. The whole algorithm can be found in Table 1.

5.1. Explicit F/B labeling

The s/t partition does not guarantee segmentation with topology constraints, hence we need to partition based on a different attribute; we propose to explicitly add an F/B label attribute to each node, which is set in the beginning of segmentation. We specify that 0 is associated to F and 1 is associated to B in the energy function (1). If there is no topology constraint, partition according to the F/B label is identical to the s/t partition described in Section 2.3, i.e., all the nodes in tree S are labeled F , and all the nodes in tree T are labeled B , thus the optimal solution of the energy function (1) is reached at the end of the computation of max-flow (by definition F is associated to tree S and B is associated to tree T). Since we are motivated by topology constraints, our algorithm must ensure that during max-flow computation, each update of the F/B label not only goes to its associated tree, i.e., to achieve lower energy of Eq. (1), but also conforms to the topology constraints. The label initialization is discussed in Section 5.2.




Fig. 4. Weight configuration for (a) edge nodes, (b) vertex node and (c) grid node (c).




Fig. 5. A default labeling of the embedding graph (a) corresponds to an optimal solution (a minimum cut of 0) of the energy function (3) without considering topology constraints. By considering the topology constraints, i.e., the foreground (red) and the background (blue) should be one connected component, the label of some door nodes and edge nodes should be changed (b,c). These changes add additional penalties to the energy function (3). (For interpretation of color mentioned in this figure the reader is referred to the web version of the article.)

Table 1

The topology cuts algorithm

1. Assign an initial label for each node
2. Set the parameters of the MRF energy function (1).
3. Compute the distance map based on the initial label map
4. Construct two trees S and T , and eight bucket priority queues:
 $\text{Act}\{S_F\}, \text{Act}\{S_B\}, \text{Act}\{T_F\}, \text{Act}\{T_B\}$
 $\text{Pas}\{S_F\}, \text{Pas}\{S_B\}, \text{Pas}\{T_F\}, \text{Pas}\{T_B\}$
5. Initialize four active sets:
 $\text{Act}\{S_F\}, \text{Act}\{S_B\}, \text{Act}\{T_F\}, \text{Act}\{T_B\}$,
 by saturating direct edges between each node to s and t
6. Topology-preserving min-cut/max-flow
 while $\text{Act}\{S_F\} \cup \text{Act}\{S_B\} \cup \text{Act}\{T_F\} \cup \text{Act}\{T_B\}$ is not empty
 - 6.1 *Inter-label maximum flow*:
 Find all non-saturated paths between the subtrees:
 $S_F \leftrightarrow T_B, S_B \leftrightarrow T_F$, until all paths are saturated
 - 6.2 *Intra-label maximum flow*:
 $\text{Act}\{S_F\} \leftarrow \text{Pas}\{S_F\}, \text{Act}\{S_B\} \leftarrow \text{Pas}\{S_B\},$
 $\text{Act}\{T_F\} \leftarrow \text{Pas}\{T_F\}, \text{Act}\{T_B\} \leftarrow \text{Pas}\{T_B\}$
 Find all non-saturated paths between the subtrees:
 $S_F \leftrightarrow T_F, S_B \leftrightarrow T_B$, until all paths are saturated

Such an explicit label attribute also resolves an inherent ambiguity in the graph cuts algorithm. When the max-flow is reached, it is possible that some nodes are isolated from both the source and sink terminals. As an example, in the image segmentation context, a node can belong either to the foreground or the background without changing the optimal solution (multi-optimal solutions). In this situation, the default assumption in [5] is: if a node does not belong to tree S , then it is assigned to tree T ¹, i.e., $T = \mathcal{V} - S$. However, a region of such isolated nodes may be inside the foreground object. By default, these nodes would be labeled background, leading to undesirable holes in the object. With our F/B label attribute, the foreground/background assignment is decoupled from the source/sink tree assignment.

Each node of the tree S and T may have one of two different labels, F or B . Thus the trees S and T are divided into four subtrees: S_F, S_B, T_F and T_B (Fig. 6).

5.2. Initialization

Next, we need to initialize the F/B label for each node to provide the topological prior information of the target segmentation. There are several ways to assign an initial label for each node, e.g., we may interactively draw some seeds to specify different connected regions and propagate them to give an initial labeling of the whole image. Alternatively, we can integrate the sampling with the segmentation, as Grabcut [27] or graph cuts for level set segmentation [6] did.




Fig. 6. Organizing the nodes by four subtrees.

As noted in [15], the concept of initialization is generally not used in the standard min-cut/max-flow algorithms, because the label of each node is not known until the min-cut is found. However, because the min-cut can not be found in a single step, any s/t cut algorithm must start from one initial state to carry on its computation. In the case of the implementation in [5], each node is initialized by saturating one of its edges to the source or sink, e.g., if its edge to the source is not saturated, then its state is set to be active and it belongs to the S tree; if both edges are saturated or there's no such edge at all, its state is set to be free. In our algorithm, we use the same technique as in [5] to initialize the active sets for each of the four subtrees S_F, S_B, T_F and T_B .

5.3. Inter/intra-label maximum flows

The computation of augmenting paths results in the change of tree membership of each node. When changes happen, we also need to update the F/B label of nodes. However, the change of tree membership does not necessarily occur along an evolving boundary between the label sets (similar to a level set evolving boundary). This makes it difficult to update the F/B label without violating the topology constraint.

If we were able to ensure that the change of tree memberships during the computation of max-flow, occurs along the boundary between the F/B labels, we would ensure that the label update is in accordance to the topology constraints. To achieve this, we give high priorities to the growing of tree nodes with different labels. That means the augmenting paths between S_F and T_B, S_B and T_F should be searched first. We call this stage *inter-label maximum flow*. After all paths between the above two pairs are saturated, there may still be non-saturated paths between nodes with the same label, namely, S_F and T_F, S_B and T_B . Thus, a second stage – *intra-label maximum flow* is employed to saturate all paths between the two subtrees pairs S_F and T_F, S_B and T_B . The details on the label updates are described in Section 6.

¹ This is based on the publicly available source code implementing [5].

5.4. Organizing the search of augmenting paths using distance maps

In the computation of max-flow, in order to localize the occurrence of the changes of tree memberships along the boundary between two label sets, those active nodes close to the boundary should be handled first in the tree growing stage. Hence, an additional attribute for each node, *DIST*, that records its distance to the boundary, is employed.

Typically, the l^k distance between two 2D points $p = (x_p, y_p)$ and $q = (x_q, y_q)$ is defined as $\|p - q\|_k = (|x_p - x_q|^k + |y_p - y_q|^k)^{\frac{1}{k}}$. Note that in the image grid, only the l^1 distance leads to integer values. The computation of the distance map with L^2 distance requires a complexity of $O(n \log n)$, whereas the computation of L^1 distance needs only $O(n)$ complexity [24]. In our algorithm, we are mostly concerned with those nodes that are closest to the boundary, which often correspond to the boundary. Here we adopt an l^1 distance map, which requires only $O(n)$ computational complexity where n is the number of nodes in the image.

In the level-set approach, the maintenance of a (signed) distance map can be costly [24]. It often requires re-initialization to make sure that the distance map is valid. Because we are only interested in those nodes that have the smallest distance values (thus they are the closest ones to the boundary), the real distance from each pixel to the boundary is not important. Thus, re-initialization is not required. The new distance value for a newly updated node can be computed as the smallest distance value among its neighboring nodes with the new label minus one.

$$DIST(n_i) = \min_{n_j, (n_i, n_j) \in \mathcal{E}, Label(n_i) \neq Label(n_j)} DIST(n_j) - 1$$

Hence, we only need to initialize the distance map one time in the beginning. The subsequent updates happen only when a node's label is changed, which requires an $O(1)$ computation for each label update.

5.5. Controlling the label propagation using the bucket priority queue data structure

At first glance, the priority queue data structure is suitable for keeping track of nodes that are closest to the boundary. However, two problems arise if we use a priority queue. First, it requires a complexity of $O(\log n)$ (n is the number of nodes in the priority queue) to extract a node, which increases the computation cost at the tree growing stage. Second and more importantly, in our topology cuts problem, those nodes that have the smallest distance value usually represent the boundary between the foreground and the background. A node on the boundary grows the tree by recruiting a new child node from its neighbors. Since the newly recruited node now has the smallest distance value, it is the first to be considered in the next stage of growing the trees by using a priority queue. This will lead to an *inhomogeneous propagation* of the boundary (evolution of the boundary from one point on the boundary as in Fig. 7 (b)), in contrast to the active contour's *homogeneous propagation* (evolution of every point on the boundary).

To reduce the cost of computation, we adopt the idea of bucket sort [9]. The range of the distance value must be within $[-m - n, m + n]$ for an $m \times n$ image with an l^1 distance map. Thus, we may allocate an array of size $2(m + n) + 1$ with each entry recording the nodes with the same distance. We also use a variable to record the current smallest distance. Because the deletion of a node can be efficiently implemented by using the node pointer, the complexity for extracting the next smallest distance node is only $O(1)$.

As for the problem of inhomogeneous propagation, we use an additional pointer to record the currently evolving boundary's distance value, which is actually a level set of the distance map. Furthermore, we restrict this pointer to not point to the entry with the smallest distance until all the nodes in the entry it currently points to have been extracted. This ensures that the boundary will evolve homogeneously (Fig. 7(c)). This makes sure that we obtain a balanced result when multiple regions (foregrounds) belong to the same tree (source or sink).

The above consideration can be efficiently implemented using the *bucket priority queue* (BPQ) data structure as first introduced in [9,11]. Fig. 8 shows the structure of the layered priority queue. Note that our new data structure can also handle other distance metric, we only need to change each entry to represent an interval instead one distance.

6. Implementation details

The overall guideline in implementing the topology cuts algorithm is to reduce the energy function while maintaining the topology constraint. Observe that if after computing the maximum flow, all nodes with label F only belong to the source tree or be free and all nodes with label B only belong to the sink tree or be free, then the energy function is minimized. Thus, during the maximum flow computation, we need to update the label of each node according to which tree it belongs to (favors).

6.1. Inter-label maximum flow

The goal of inter-label maximum flow is to quickly evolve the boundary between the F -labeled and B -labeled regions under topology constraints. In this stage, all the augmenting paths between the two subtree pairs, (S_F, T_B) and (S_B, T_F) , are searched. In each search step, we either find an augmenting path or grow these subtrees by setting each active node's neighbors with non-saturated edges as its new children, if the conditions stated below are met. The distance value and the label of a new child q are updated if q 's original label is different from that of its parent.

The conditions for recruiting a new neighbor q by active node p are:

- (1) If the label of q is the same as that of p , then q is recruited if and only if q is free.
- (2) If q has a different label and q is a simple point, then q is recruited if either of the following conditions is met:




Fig. 7. A synthetic example illustrating the importance of using our BPQ data structure in organizing the search of the maximum flow.




Fig. 8. The bucket priority queue data structure.

- (a) q is free, or
- (b) p is associated to the tree that q belongs to (Section 5.1).

The above conditions ensure that the F/B labels are updated with respect to the topological constraints while searching for the max-flow.

6.2. Intra-label maximum flow

The goal of the intra-label maximum flow is to saturate all the single-label paths between tree S and tree T , and change the label of any node for which the following conditions are met.

The label of a node p is changed only when (1) p 's opposite label is associated to the tree that p belongs to (Section 5.1), and (2) p is a simple point. Note that because we do not know if p 's tree mem-

bership will change in the subsequent computations, we should not change the label of p until the end of this stage. Thus, we simply record all the nodes that meet the above two conditions as candidate nodes and finalize label changes after the computation of max-flow in this stage. The labels of these candidate nodes that still satisfy the above conditions will be changed and inserted into their corresponding active sets for the inter-label maximum flow stage in the next iteration.

7. Analysis of the topology cuts algorithm

7.1. Convergence

Convergence of our algorithm is ensured by (1) all augmenting paths between s and t are guaranteed to be saturated, and (2) the algorithm will stop within two iterations.




Fig. 9. An example of segmentation of knee image illustrates the advantage of our algorithm over graph cuts. Topology-free segmentation is usually not desirable for medical applications. The second row shows the corresponding label maps of the results in the first row.




Fig. 10. An example of segmentation of ventricular image. The second row shows the corresponding label maps of the results in the first row.




Fig. 11. Comparison between our algorithm and level sets [14] results. The result by level sets is similar to that of [33] except that we applied a different initialization.

To verify the first claim, observe that in the first iteration, all augmenting paths between the underlying four subtrees with two different label sets are saturated after finding the inter-label maximum flow. Likewise, between the stages of searching the intra-label maximum flow and changing labels, all augmenting paths with the same labels are also saturated. The only non-saturated paths left are those between the current S and T trees with different label sets. Then at least one of these nodes along such a path should reside on the boundary. According to our intra-label flow algorithm, this node must be recorded and the path will be found and saturated in the next iteration.

For the second claim, note that the remaining non-saturated paths are those crossing two different label sets at the end of the first iteration. According to our inter-label maximum flow implementation, all such paths must be saturated in the next iteration. There does not exist any other non-saturated path after the second iteration and hence the whole algorithm must stop within two iterations.

Note that without considering the topology constraint (updating the labels without considering the simplicity of the nodes), our algorithm is actually another implementation of the min-cut/max-flow algorithm.

7.2. Topology preservation

This can be easily verified by looking into our algorithm to see that, before each update of the label, the simple point condition is always checked, i.e., the label of a node is changed

only if its change does not affect the global topology of the image.

7.3. Time complexity

The time complexity can be shown by comparing the differences between our algorithm and the implementation of graph cuts in [5]. Both algorithms search the augmenting paths by growing two trees from s and t , respectively. The main difference here is that we divide the search into two stages. Since our algorithm stops within two iterations, each node is traversed at most four times (one time for each stage). The bucket priority queue data structure ensures that the selection of active nodes needs an $O(1)$ operation. The update of the distance value for a node also needs an $O(1)$ operation. In addition, an initialization of the distance map is computed in the beginning, which requires an $O(n)$ operation where n is the number of nodes (pix-

Table 2

Performance comparison among graph cuts, topology cuts and level sets with the same initialization and parameter values

Image	Size	Graph cuts (ms)	Topology cuts (ms)	Level sets (s)
Synthetic	250×180	16	16	48
Brain	200×200	20	31	55
Ventricle	256×256	16	47	213
Knee	400×400	63	109	359

For level sets result, the iteration stops when the average update of each pixel is less than 0.001.




Fig. 12. Example of using topology cuts algorithm for interactive object cutout (without border matting).




Fig. 13. Results for topology-preserving hand tracking (also see our video).

els) in the image. In total, our algorithm only adds a constant factor to the complexity of the original algorithm in the worst case. In practice, our algorithm works sufficiently fast since the number of active nodes is significantly reduced in the second iteration.

8. Experimental results

We apply our topology cuts algorithm to two problems: image segmentation and object tracking.² All results were run on a PC equipped with an Intel Pentium M 2.0 GHz processor and 1.5 G memory.

8.1. Results for image segmentation

To verify our algorithm, we use the discrete piecewise Mumford–Shah style energy function [10]. The Mumford–Shah model is widely applied in the level set literature and it works well for grayscale images without much texture part, such as CT medical images. It also allows us to use a level set style initialization that integrates topology initialization and seed assignment. The parameters of the MRF energy function (1) are defined as follows:

$$\begin{aligned} D(x_p) &= \lambda((u_p - c^B)^2 x_p + (u_p - c^F)^2 (1 - x_p)) \\ V_{pq}(x_p, x_q) &= x_p(1 - x_q) + (1 - x_p)x_q \end{aligned} \quad (4)$$

Here, u_p denotes the gray value of the image at p , c^F and c^B denote the mean gray value of the pixels with the label F and B , respectively (the initial labeling is assigned by users). We use $\lambda = 10$ in our experiments, and solve the energy function in one step instead of iteratively estimating the mean gray value of the foreground and background.

Traditionally, graph cuts algorithms only take account of color and coherence between neighboring pixels. Whereas, solely exploiting the color similarity can not fully guarantee meaningful segmentation results for medical applications. An immediate example is demonstrated in Fig. 9. Our algorithm gives a result (Fig. 9(c)) that faithfully conforms to the initialization, which is more meaningful than that of the standard graph cuts algorithm (Fig. 9(b)). Fig. 10 shows another example. Fig. 11 shows an example of comparison between our algorithm and the level sets [14] algorithm. As pointed out in [33], the evolution of the curve by Han's method [14] could become “stuck” when two boundaries meet in order to respect the topological constraint (Fig. 11(b)). Our algorithm achieves visually better results because of the label update rules explained in Section 5.5. The main advantage of our

graph based algorithm over the level sets method is the speedup. Furthermore our algorithm is based on the discrete graphical MRF models which are often considered more general than the continuous PDE based models since higher order or longer distance relations among pixels can be explored. Table 2 shows the comparison in speed between level sets and our topology cuts algorithm. It also shows the comparison between the graph cuts implementation [5] and our algorithm.

8.2. Results for interactive object cutout

Natural images often contain richly textured parts. Modeling them using the pairwise MRF model is insufficient since it only considers local interactions of the pixels. Moreover, its result tends to be sensitive to the choice of parameters, i.e., the weights balancing the data term (the first term of (1)) and the smoothness term (the second term of (1)) are often difficult to determine. With a large weight for the smoothness term, a large part of the background may be segmented. And with a large weight for the data term, holes and outliers may easily be generated in the segmented object. To account for this, some global properties should be introduced to appropriately model the natural image, such as global topology information which is often available in many applications.

Fig. 12 illustrates the result of applying our topology cuts algorithm for interactive object cutout. We use small values for the weight that balances the two terms, e.g., $\gamma = 10$. From Fig. 12(b) and (c) we can see that holes and small outliers are generated inside the object by using the standard graph cuts algorithm. Instead, by applying our topology cuts algorithm, we obtain a complete result as shown in Fig. 12(d) and (e). Our method might still suffer if outliers are large and conform with the topology prior. However, our algorithm significantly reduces sensitivity to weight selection.

8.3. Results for object tracking

Our topology cuts algorithm can be applied to topology-preserving object tracking. The user assigns the initial contours (or they can be automatically located) containing the topology prior information for the first frame. For all the subsequent frames, the segmentation result from the previous frame is used as the initialization before the topology cuts algorithm is applied, thus the topology information is propagated from the first frame to the other frames. Fig. 13 illustrates the results of a simple implementation for a hand tracking case. Topology is correctly preserved as shown in Fig. 13(b),(e),(f), even when the two hands meet.³

² More results and comparisons are submitted as supplementary materials.

³ The tracked video sequence is in the supplemental materials.

9. Conclusions and future work

We proposed a new algorithm for solving a subset of MRF functions that can be addressed by graph cuts while respecting topological constraints. It combines certain advantages of level sets and graph cuts. The idea of boundary evolution is introduced into the graph cuts framework by using the explicit F/B label attribute. Rather than evolving the boundary in a gradient descent manner to update the distance function as level sets do, the boundary evolution of the F/B label set is driven by the computation of max-flow, which is fast and stable. The bucket priority queue data structure ensures that there is no increase in computational complexity compared with the existing graph cuts algorithms.

In the near future, we plan to extend the topology cuts algorithm to allow soft constraints. We would also like to apply our new algorithm to other vision problems such as stereo, 3D reconstruction. One promising direction is the incorporation of other prior knowledge into the min-cut/max-flow algorithm, e.g., the curvature of the boundary can be approximately encoded into the distance function. With our framework we would be able to design a new cut algorithm that considers the smoothness of the boundary as well.

Acknowledgements

We are grateful to Prof. Y. Boykov for enlightening discussions on this work. This work was partially supported by NIDA Grant: 1 R01 DA020949-01, NSF Grants: CNS-0627645, CNS-0627645, IIS-0527585 and 863 Program of China (No. 2006AA01Z314).

References

- [1] R.K. Ahuja, T.L. Magnanti, J.B. Orlin, *Network Flows: Theory, Algorithms, and Applications*, Prentice Hall, 1993.
- [2] G. Bertrand, Simple points, topological numbers and geodesic neighborhoods in cubic grids, *Pattern Recognition Letters* 15 (10) (1994) 1003–1011.
- [3] Y. Boykov, G. Funka-Lea, Graph cuts and efficient N-D image segmentation, *IJCV* 70 (2) (2006) 109–131.
- [4] Y. Boykov, V. Kolmogorov, Computing geodesics and minimal surfaces via graph cuts, in: *ICCV* 03, 2003, pp. 26–33.
- [5] Y. Boykov, V. Kolmogorov, An experimental comparison of min-cut/max-flow algorithms for energy minimization in vision, *PAMI* 26 (9) (2004) 1124–1137.
- [6] Y. Boykov, V. Kolmogorov, D. Cremers, A. Delong, An integral solution to surface evolution PDEs via geo-cuts, in: *ECCV* 06, 2006, pp. 409–422.
- [7] Y. Boykov, O. Veksler, Graph cuts in vision and graphics: theories and applications, in: N. Paragios, Y. Chen, O. Faugeras (Eds.), *The Handbook of Mathematical Models in Computer Vision*, Springer-Verlag, 2006, pp. 79–96.
- [8] Y. Boykov, O. Veksler, R. Zabih, Fast approximate energy minimization via graph cuts, *PAMI* 23 (11) (2001) 1222–1239.
- [9] R. Brown, Calendar queues: a fast $O(1)$ priority queue implementation for the simulation event set problem, *Commun. ACM* 31 (10) (1998) 1220–1227.
- [10] T. Chan, L. Vese, Active contour model without edges, *IEEE Transactions on Image Processing* 10 (2) (2001) 266–277.
- [11] B.V. Cherkassky, A.V. Goldberg, On implementing the push-relabel method for the maximum flow problem, *Algorithmica* 19 (4) (1997) 390–410.
- [12] D. Freedman, T. Zhang, Interactive graph cut based segmentation with shape priors, in: *CVPR* 05, 2005, pp. 755–762.
- [13] S. German, D. German, *Stochastic Relaxation, Gibbs Distributions, and the Bayesian Restoration of Images*, MIT Press, Cambridge, MA, USA, 1988.
- [14] X. Han, C. Xu, J.L. Prince, A topology preserving level set method for geometric deformable models, *PAMI* 25 (6) (2003) 755–768.
- [15] O. Juan, Y. Boykov, Active graph cuts, in: *CVPR* 06, 2006, pp. 1023–1029.
- [16] N. Komodakis, G. Tziritas, A new framework for approximate labeling via graph cuts, in: *ICCV* 05, 2005, pp. 1018–1025.
- [17] N. Komodakis, Nikos Paragios, Georgios Tziritas, MRF optimization via dual decomposition: message-passing revisited, in: *ICCV* 07.
- [18] N. Komodakis, G. Tziritas, N. Paragios, Fast, approximately optimal solutions for single and dynamic MRFs, in: *CVPR* 07, 2007, pp. 1–8.
- [19] V. Kolmogorov, Y. Boykov, What metrics can be approximated by geo-cuts, or global optimization of length/area and flux, in: *ICCV* 05, 2005, pp. 564–571.
- [20] V. Kolmogorov, R. Zabih, What energy functions can be minimized via graph cuts?, *PAMI* 26 (2) (2004) 147–159.
- [22] M.P. Kumar, P.H.S. Torr, A. Zisserman, OBJ CUT, in: *CVPR* 05, 2005, pp. 18–25.
- [23] H. Lombaert, Y. Sun, L. Grady, C. Xu, A multilevel banded graph cuts method for fast image segmentation, in: *ICCV* 05, vol. 1, 2005, pp. 259–265.
- [24] S. Osher, R. Fedkiw (Eds.), *Level Set Methods and Dynamic Implicit Surfaces*, Springer-Verlag, 1998.
- [25] S. Osher, J. Sethian, Fronts propagating with curvature-dependent speed: Algorithms based on the Hamilton-Jacobi formulation, *J. Comput. Phys.* 79 (2) (1988) 12–49.
- [26] N. Paragios, Y. Chen, O. Faugeras (Eds.), *The Handbook of Mathematical Models in Computer Vision*, Springer-Verlag, 2005.
- [27] C. Rother, V. Kolmogorov, A. Blake, Grabcut: interactive foreground extraction using iterated graph cuts, *ACM Trans. Graph.* 23 (3) (2004) 309–314.
- [28] W. Shi, C. Su, The rectilinear steiner arborescence problem is np-complete, in: *SODA* 00, 2000, pp. 780–787.
- [29] Y. Shi, W.C. Karl, Real-time tracking using level sets, in: *CVPR* 05, 2005, pp. 34–41.
- [30] V.V. Vazirani, *Approximation Algorithms*, Springer-Verlag, 2001.
- [31] M.R. Garey, D.S. Johnson, The rectilinear steiner tree problem is NP-complete, *SIAM J. Appl. Math.* 32 (4) (1977) 826–834.
- [32] Florent Ségonne, *Segmentation of Medical Images under Topological Constraints*, PhD Thesis, MIT, 2005.
- [33] G. Sundaramoorthi, A. Yezzi, More-than-topology-preserving flows for active contours and polygons, in: *ICCV* 05, vol. 2, 2005, pp. 1276–1283.